Table of Contents

Preface ..... **ix**

1. Why Machine Learning and Security? **1**
Cyber Threat Landscape 3
The Cyber Attacker's Economy 7
A Marketplace for Hacking Skills 7
Indirect Monetization 8
The Upshot 8
What Is Machine Learning? 9
What Machine Learning Is Not 10
Adversaries Using Machine Learning 11
Real-World Uses of Machine Learning in Security 12
Spam Fighting: An Iterative Approach 14
Limitations of Machine Learning in Security 23

2. Classifying and Clustering **25**
Machine Learning: Problems and Approaches 25
Machine Learning in Practice: A Worked Example 27
Training Algorithms to Learn 32
Model Families 33
Loss Functions 35
Optimization 36
Supervised Classification Algorithms 40
Logistic Regression 40
Decision Trees 42
Decision Forests 45
Support Vector Machines 47
Naive Bayes 49

k-Nearest Neighbors 51
Neural Networks 53
Practical Considerations in Classification 55
Selecting a Model Family 55
Training Data Construction 56
Feature Selection 59
Overfitting and Underfitting 60
Choosing Thresholds and Comparing Models 62
Clustering 65
Clustering Algorithms 65
Evaluating Clustering Results 75
Conclusion 77

3. Anomaly Detection. 79
When to Use Anomaly Detection Versus Supervised Learning 80
Intrusion Detection with Heuristics 81
Data-Driven Methods 82
Feature Engineering for Anomaly Detection 85
Host Intrusion Detection 85
Network Intrusion Detection 89
Web Application Intrusion Detection 92
In Summary 93
Anomaly Detection with Data and Algorithms 93
Forecasting (Supervised Machine Learning) 95
Statistical Metrics 106
Goodness-of-Fit 107
Unsupervised Machine Learning Algorithms 112
Density-Based Methods 116
In Summary 118
Challenges of Using Machine Learning in Anomaly Detection 119
Response and Mitigation 120
Practical System Design Concerns 121
Optimizing for Explainability 121
Maintainability of Anomaly Detection Systems 123
Integrating Human Feedback 123
Mitigating Adversarial Effects 123
Conclusion 124

4. Malware Analysis. 125
Understanding Malware 126
Defining Malware Classification 128
Malware: Behind the Scenes 131

Feature Generation 145
Data Collection 146
Generating Features 147
Feature Selection 171
From Features to Classification 174
How to Get Malware Samples and Labels 178
Conclusion 179

5. Network Traffic Analysis. 181
Theory of Network Defense 183
Access Control and Authentication 183
Intrusion Detection 184
Detecting In-Network Attackers 185
Data-Centric Security 185
Honeypots 186
Summary 186
Machine Learning and Network Security 187
From Captures to Features 187
Threats in the Network 193
Botnets and You 196
Building a Predictive Model to Classify Network Attacks 203
Exploring the Data 205
Data Preparation 210
Classification 215
Supervised Learning 216
Semi-Supervised Learning 222
Unsupervised Learning 223
Advanced Ensembling 229
Conclusion 233

6. Protecting the Consumer Web. 235
Monetizing the Consumer Web 236
Types of Abuse and the Data That Can Stop Them 237
Authentication and Account Takeover 237
Account Creation 244
Financial Fraud 248
Bot Activity 251
Supervised Learning for Abuse Problems 256
Labeling Data 256
Cold Start Versus Warm Start 257
False Positives and False Negatives 258
Multiple Responses 259

Large Attacks 259
Clustering Abuse 259
Example: Clustering Spam Domains 261
Generating Clusters 262
Scoring Clusters 266
Further Directions in Clustering 271
Conclusion 272

7. Production Systems. 275
Defining Machine Learning System Maturity and Scalability 275
What's Important for Security Machine Learning Systems? 277
Data Quality 277
Problem: Bias in Datasets 277
Problem: Label Inaccuracy 279
Solutions: Data Quality 279
Problem: Missing Data 280
Solutions: Missing Data 281
Model Quality 284
Problem: Hyperparameter Optimization 285
Solutions: Hyperparameter Optimization 285
Feature: Feedback Loops, A/B Testing of Models 290
Feature: Repeatable and Explainable Results 293
Performance 297
Goal: Low Latency, High Scalability 297
Performance Optimization 298
Horizontal Scaling with Distributed Computing Frameworks 300
Using Cloud Services 305
Maintainability 307
Problem: Checkpointing, Versioning, and Deploying Models 307
Goal: Graceful Degradation 308
Goal: Easily Tunable and Configurable 309
Monitoring and Alerting 310
Security and Reliability 311
Feature: Robustness in Adversarial Contexts 312
Feature: Data Privacy Safeguards and Guarantees 312
Feedback and Usability 313
Conclusion 314

8. Adversarial Machine Learning. 315
Terminology 316
Importance of Adversarial ML 317
Security Vulnerabilities in Machine Learning Algorithms 318

Attack Transferability 320
Attack Technique: Model Poisoning 322
Example: Binary Classifier Poisoning Attack 325
Attacker Knowledge 331
Defense Against Poisoning Attacks 332
Attack Technique: Evasion attack 334
Example: Binary Classifier Evasion Attack 335
Defense Against Evasion Attacks 340
Conclusion 341

A. Chapter 2 Appendix. 343

B. Chapter 5 Appendix. 351

Index. 355